Data Science Career Program Assignments (Theoretical & Practical)

Module 1: Introduction to Data Science

Theoretical Assignments:

1. **Report on the Evolution of DataScience** – Trace how data science has evolved over the years.
2. **Presentation: Applications of Data Science in Various Industries** –Students present real-world use cases in Healthcare, Finance, Retail, etc.
3. **Essay: How Generative AI is Transforming DataScience –**Explain the impact of AI- generated data.

PracticalTasks:

Task1: CaseStudy Analysis on Data-Driven DecisionMaking

- Research a company (e.g.,Netflix,Amazon) and analyze how they use data for decision-making.
- Identify the problem ,data collected,techniques used, and impact on business. ·

Deliverable : A report with insights and a possible presentation.

Task2:Identifying Real-World DataScience Problems

- Choose an industry (Healthcare, Finance, E-commerce, etc.).

- List five potential problems where data science can provide solutions
- **Deliverable**: A document explaining each problem and possible data-driven solutions

Module2: Introduction to Python

Theoretical Assignments:

1. Explain Python's Role in Data Science
2. Presentation: Object-Oriented Programming (OOP) Concepts in Python

Practical Tasks:

Task1 :Python Fundamentals

1. Write a python program to sum of the first n positive integers.
2. Write a Python program to count occurrences of a substring in a string.
3. Write a Python program to count the occurrences of each word in a given sentence.
4. Write a Python program to get a single string from two given strings, separated by a space and swap the first two characters of each string.
5. Write a Python program to add 'ing' at the end of a given string (length should be at least 3). If the given string already ends with 'ing' then add 'ly' instead If the string length of the given string is less than 3, leave it unchanged
6. Write a Python program to find the first appearance of the substring 'not' and 'poor' from a given string, if 'not' follows the 'poor', replace the whole 'not'...'poor' substring with 'good'. Return the resulting string
7. Program to find Greatest Common Divisor of two numbers. For example, the GCD of 20 and 28 is 4 and the GCD of 98 and 56 is 14.
8. Write a Python program to check whether a list contains a sublist.
9. Write a Python program to find the second smallest number in a list.
10. Write a Python program to get unique values from a list.
11. Write a Python program to unzip a list of tuples into individual lists.
12. Write a Python program to convert a list of tuples into a dictionary
13. Write a Python program to sort a dictionary (ascending /descending) by value.
14. Write a Python program to find the highest 3 values in a dictionary.

15. Given a number n, write a python program to make and print the list of Fibonacci series up to n. Input : n=7 Hint : first 7 numbers in the series Expected output :

    First few Fibonacci numbers are 0, 1, 1, 2, 3, 5, 8, 13

16. Counting the frequencies in a list using a dictionary in Python. Input : [1, 1, 1, 5, 5, 3, 1, 3, 3, 1,4, 4, 4, 2, 2, 2, 2]

    Expected output : 1 : 5 , 2 : 4 , 3 : 3 , 4 : 3 , 5 : 2

17. Write a python program using function to find the sum of odd series and even series

    Odd series: 12/ 1! +32/ 3! + 52/ 5!+……n

    Even series: 22/ 2! + 42/ 4! + 62/ 6!+……n

18. Python Program to Find Factorial of Number Using Recursion
19. Write a Python function that takes a list and returns a new list with unique elements of the first list.
20. Mini project :

    Problem Statement : Password Generator

    Make a program to generate a strong password using the input given by the user. To generate a password, randomly take some words from the user input and then include numbers, special characters and capital letters to generate the password. Also, keep a check that password length is more than 8 characters.

    Note: Include Exception handling wherever required. Also, make a ‘User’ class and store the details like user id, name and password of each user as a tuple.

MODULE-3: Working with Databases

Theoretical Assignments:

1. Compare SQL and NoSQL Databases
2. Presentation on Normalization vs. Denormalization

Dataset: https://github.com/TopsCode/Data-Science

Practical Tasks for SQL:

1. write a SQL query to find customers who are either from the city 'New York' or who do not have a grade greater than 100. Return customer_id, cust_name, city, grade, and salesman_id.
2. write a SQL query to find all the customers in 'New York' city who have a grade value above 100. Return customer_id, cust_name, city, grade, and salesman_id.
3. Write a SQL query that displays order number, purchase amount, and the achieved and unachieved percentage (%) for those orders that exceed 50% of the target value of 6000.
4. write a SQL query to calculate the total purchase amount of all orders. Return total purchase amount.
5. write a SQL query to find the highest purchase amount ordered by each customer. Return customer ID, maximum purchase amount.
6. write a SQL query to calculate the average product price. Return average product price.
7. write a SQL query to find those employees whose department is located at 'Toronto'. Return first name, last name, employee ID, job ID.
8. write a SQL query to find those employees whose salary is lower than that of employees whose job title is "MK_MAN". Exclude employees of the Job title 'MK_MAN'. Return employee ID, first name, last name, job ID.
9. write a SQL query to find all those employees who work in department ID 80 or 40. Return first name, last name, department number and department name.
10. write a SQL query to calculate the average salary, the number of employees receiving commissions in that department. Return department name, average salary and number of employees.
11. write a SQL query to find out which employees have the same designation as the employee whose ID is 169. Return first name, last name, department ID and job ID.

12. write a SQL query to find those employees who earn more than the average salary. Return employee ID, first name, last name.
13. write a SQL query to find all those employees who work in the Finance department. Return department ID, name (first), job ID and department name.
14. From the following table, write a SQL query to find the employees who earn less than the employee of ID 182. Return first name, last name and salary.
15. Create a stored procedure CountEmployeesByDept that returns the number of employees in each department
16. Create a stored procedure AddNewEmployee that adds a new employee to the database.
17. Create a stored procedure DeleteEmployeesByDept that removes all employees from a specific department
18. Create a stored procedure GetTopPaidEmployees that retrieves the highest-paid employee in each department.
19. Create a stored procedure PromoteEmployee that increases an employee's salary and changes their job role.
20. Create a stored procedure AssignManagerToDepartment that assigns a new manager to all employees in a specific department.

Practical Tasks for NoSQL:

1. Retrieve all employee records
2. Find employees who work in the IT department
3. Find employees who have a salary greater than 70,000
4. Find employees who joined after 2018
5. Find employees between the ages of 30 and 40
6. Increase the salary of all employees in the Finance department by 5%
7. Delete employees who joined before 2010
8. Find the highest-paid employee

Module 4: Essential Statistics & Mathematics for Data Science

Theoretical Assignments:

1. Explain Key Statistical Concepts
2. Presentation: Importance of Probability in Machine Learning

Practical Tasks

1. Business Problem: A car rental company wants to analyze the rental durations of its customers to understand the typical rental period and optimize its pricing and fleet management strategies. Data: Let's consider the rental durations (in days) for a sample of 50 customers:

   3, 2, 5, 4, 7, 2, 3, 3, 1, 6, 4, 2, 3, 5, 2, 4, 2, 1, 3, 5, 6, 3, 2, 1, 4, 2, 4, 5, 3, 2, 7, 2, 3, 4, 5, 1, 6, 2, 4, 3, 5, 3, 2, 4, 2, 6, 3, 2, 4, 5

   Question:

   a. Mean: What is the average rental duration for customers at the car rental company?
   b. Median: What is the typical or central rental duration experienced by customers?
   c. Mode: Are there any recurring or most frequently occurring rental durations for customers?

   By answering these questions using the mean, median, and mode, the car rental company can gain insights into the average rental duration, understand the most common rental periods, and make informed decisions regarding pricing, fleet size, and availability. Additionally, this analysis can help the company optimize resource allocation, plan for peak demand periods, and enhance customer satisfaction by aligning service offerings with customers' typical rental needs.

2. An e-commerce platform wants to analyze the delivery times of its shipments to understand the variability in order fulfillment and optimize its logistics operations.

   Data:

   Let's consider the delivery times (in days) for a sample of 50 shipments:

   3, 5, 2, 4, 6, 2, 3, 4, 2, 5,

   7, 2, 3, 4, 2, 4, 2, 3, 5, 6,

   3, 2, 1, 4, 2, 4, 5, 3, 2, 7,

   2, 3, 4, 5, 1, 6, 2, 4, 3, 5,

   3, 2, 4, 2, 6, 3, 2, 4, 5, 3

   Questions:

   a. Range: What is the range of the delivery times?
   b. Variance: What is the variance of the delivery times?
   c. Standard Deviation: What is the standard deviation of the delivery times?

   By answering these questions using different measures of dispersion, the e-commerce platform can gain insights into the variability in delivery times, identify any bottlenecks in the logistics process, and make informed decisions regarding shipment tracking,customer expectations, and service level agreements.

3. A company wants to analyze the monthly revenue generated by one of its products to understand its performance and variability.

   Data: Let's consider the monthly revenue (in thousands of dollars) for the past 12 months:

   $120, $150, $110, $135, $125, $140, $130, $155, $115, $145, $135, $130

   Questions:

   a. Measure of Central Tendency: What is the average monthly revenue for the product?
   b. Measure of Dispersion: What is the range of monthly revenue for the product?

By answering these questions, the company can gain insights into the average revenue generated by the product and understand the range or variability in its monthly revenue, which can help with financial planning, forecasting, and evaluating the product's performance.

4. A transportation company wants to analyze the fuel efficiency of its vehicle fleet to identify any variations across different vehicle models.

   Data: Let's consider the fuel efficiency (in miles per gallon, mpg) for a sample of 50 vehicles:

   Model A: 30, 32, 33, 28, 31, 30, 29, 30, 32, 31,

   Model B: 25, 27, 26, 23, 28, 24, 26, 25, 27, 28,

   Model C: 22, 23, 20, 25, 21, 24, 23, 22, 25, 24,

   Model D: 18, 17, 19, 20, 21, 18, 19, 17, 20, 19,

   Model E: 35, 36, 34, 35, 33, 34, 32, 33, 36, 34

   Questions:

   a. Measure of Central Tendency: What is the average fuel efficiency for each vehicle model?
   b. Measure of Dispersion: What is the range of fuel efficiency for each vehicle model?
   c. Measure of Dispersion: What is the variance of the fuel efficiency for each vehicle model?

   By answering these questions, the transportation company can gain insights into the average fuel efficiency of different vehicle models, understand the variations or spread in the fuel efficiency, and make informed decisions regarding fleet management, vehicle selection, and fuel consumption optimization.

5. A manufacturing company wants to analyze the defect rates of its production line to identify the frequency of different types of defects.

   Data: Let's consider the types of defects and their corresponding frequencies observed in a sample of 200 products:

Defect Type: A, B, C, D, E, F, G

Frequency: 30, 40, 20, 10, 45, 25, 30

Questions:

a. Bar Chart: Create a bar chart to visualize the frequency of different defect types.
b. Most Common Defect: Which defect type has the highest frequency?
c. Histogram: Create a histogram to represent the defect frequencies.

By answering these questions using a bar chart and histogram, the manufacturing company can visually understand the distribution of defect types, identify the most common defect, and prioritize quality control efforts to address the prevalent issues.

6. A survey was conducted to analyze the satisfaction ratings of customers on a scale of 1 to 5 for a specific product.

Data: Let's consider the satisfaction ratings from 200 customers:

Ratings:

4, 5, 3, 4, 4, 3, 2, 5, 4, 3, 5, 4, 2, 3, 4, 5, 3, 4, 5, 3, 4, 3, 2, 4, 5, 3, 4, 5, 4, 3, 3, 4, 5, 2, 3, 4, 4, 3, 5, 4, 3, 4, 5, 4, 2, 3, 4, 5, 3, 4, 5, 4, 3, 4, 5, 3, 4, 5, 4, 3, 3, 4, 5, 2, 3, 4, 4, 3, 5, 4, 3, 4, 5, 4, 2, 3, 4, 5, 3, 4, 5, 4, 3, 4, 5, 3, 4, 5, 4, 3, 3, 4, 5, 2, 3, 4, 4, 3, 5, 4

Questions:

a. Skewness: Calculate the skewness of the satisfaction ratings.
b. Kurtosis: Calculate the kurtosis of the satisfaction ratings.
c. Interpretation: Based on the skewness and kurtosis values, what can be inferred about the satisfaction ratings distribution?

By answering these questions using measures of skewness and kurtosis, the survey can assess the skewness and peakedness of the satisfaction ratings, determine if the ratings are skewed towards positive or negative evaluations, and understand the distribution characteristics of customer satisfaction.

7. A research study wants to analyze the weight distribution of a sample of individuals to assess their health and body composition.

Data: Let's consider the weights (in kilograms) of a sample of 100 individuals: Weights:

55, 60, 62, 65, 68, 70, 72, 75, 78, 80, 82, 85, 88, 90, 92, 95, 100, 105, 110, 115, 120, 125, 130, 135, 140, 145, 150, 155, 160, 165, 170, 175, 180, 185, 190, 195, 200, 205, 210, 215, 220, 225, 230, 235, 240, 245, 250, 255, 260, 265, 270, 275, 280, 285, 290, 295, 300, 305, 310, 315, 320, 325, 330, 335, 340, 345, 350, 355, 360, 365, 370, 375, 380, 385, 390, 395, 400, 405, 410, 415, 420, 425, 430, 435, 440, 445, 450, 455, 460, 465, 470, 475, 480, 485, 490, 495, 500, 505, 510, 515

Questions:

a. Quartiles: Calculate the first quartile (Q1), median (Q2), and third quartile (Q3) of the weight distribution.
b. Percentiles: Calculate the 15th percentile, 50th percentile, and 85th percentile of the weight distribution.
c. Interpretation: Based on the quartiles and percentiles, what can be inferred about the weight distribution of the individuals?

By answering these questions using quartiles and percentiles, the research study can understand the weight distribution and identify the weight ranges at different percentiles, such as underweight, normal weight, overweight, and obese categories. This information

8. A researcher wants to examine the relationship between the hours spent studying and the exam scores of a group of students.

Data: Let's consider the number of hours spent studying and the corresponding exam scores for a sample of 30 students:

Hours Spent Studying: 10, 12, 15, 18, 20, 22, 25, 28, 30, 32, 35, 38, 40, 42, 45, 48, 50, 52, 55, 58, 60, 62, 65, 68, 70, 72, 75, 78, 80, 82 Exam Scores: 60, 65, 70, 75, 80, 82, 85, 88, 90, 92, 93, 95, 96, 97, 98, 99, 100, 102, 105, 106, 107, 108, 110, 112, 114, 115, 116, 118, 120, 122

Question:

Calculate the correlation coefficient between the hours spent studying and the exam scores. Interpret the value of the correlation coefficient and explain the nature of the relationship between studying hours and exam scores.

By analyzing the correlation coefficient, the researcher can determine the strength and direction of the relationship between studying hours and exam scores. This information can provide insights into the effectiveness of studying and help students and educators make informed decisions about study habits and academic performance.

9. A multiple-choice test consists of 10 questions, each with four possible answers. If a student randomly guesses on each question, what is the probability of getting at least 8 questions correct? Data: Number of questions (n) = 10, Number of possible answers per question (k) = 4
10. The lifetimes of a certain brand of light bulbs are normally distributed with a mean of 1000 hours and a standard deviation of 100 hours. What is the probability that a randomly selected light bulb lasts between 900 and 1100 hours? Data: Mean lifetime (μ) = 1000 hours, Standard deviation (σ) = 100 hours, Lifetime range (lower limit x1, upper limit x2)
11. A study is conducted to estimate the mean height of a population. A random sample of 100 individuals is selected, and their heights are measured. Calculate a 95% confidence interval for the population mean height, given that the sample mean height is 170 cm and the sample standard deviation is 8 cm. Data: Sample size (n) = 100, Sample mean ($\bar{x}$) = 170 cm, Sample standard deviation (s) = 8 cm, Confidence level = 95%

    Explanation: In this problem, we use a sample to estimate the population mean height. By calculating a confidence interval, we provide a range of plausible values for the population mean. The 95% confidence level indicates that we are 95% confident that the true population mean height falls within the calculated interval.

12. A manufacturing company claims that the average weight of its product is 500 grams. To test this claim, a random sample of 25 products is selected, and their weights are measured. The sample mean weight is found to be 510 grams with a sample standard deviation of 20 grams. Perform a hypothesis test to determine if there is evidence to support the company's claim. Data: Sample size (n) = 25, Sample mean ($\bar{x}$) = 510 grams, Sample standard deviation (s) = 20 grams, Population mean (μ) = 500 grams

Explanation: In this problem, we are conducting a hypothesis test to assess whether the sample mean weight provides evidence to support the company's claim about the population mean weight. The null hypothesis (H0) assumes that the population mean weight is equal to the claimed value, while the alternative hypothesis (Ha) suggests otherwise.

13. A researcher wants to test whether a new teaching method improves student performance. A random sample of 50 students is divided into two groups: one group taught using the new method and the other using the traditional method. The average test scores of the two groups are compared. State the null and alternative hypotheses for this study. Data: Sample size (n) = 50, Test scores of the two groups

    Explanation: In this problem, we are interested in comparing the means of two groups (new method vs. traditional method). The null hypothesis (H0) states that there is no significant difference between the means, while the alternative hypothesis (Ha) suggests that there is a significant difference

14. A survey is conducted to estimate the proportion of people in a city who support a particular policy. A random sample of 500 individuals is surveyed, and 320 of them express support for the policy. Calculate a 90% confidence interval for the population proportion, given the sample proportion. Data: Sample size (n) = 500, Number of successes (x) = 320, Confidence level = 90%

    Explanation: In this problem, we aim to estimate the population proportion based on the sample proportion. By constructing a confidence interval, we provide a range of plausible values for the population proportion. The 90% confidence level indicates that we are 90% confident that the true population proportion falls within the calculated interval.

15. A company sells smartphones, and the number of defects per batch follows a Poisson distribution with a mean of 2 defects. What is the probability of having exactly 3 defects in a randomly selected batch? Data: Mean number of defects (λ) = 2, Number of defects (x) = 3

    Explanation: The problem involves a discrete distribution (Poisson) because we are dealing with the count of defects in a batch of smartphones. The Poisson

distribution models the probability of a given number of events occurring within a fixed interval of time or space.

Module 5:Data Collection, Cleaning, Visualization, and Analysis

Theoretical Assignments:

1. Write an essay: Importance of Data Cleaning in Data Science
2. Presentation: Data Visualization Techniques & Best Practices

Practical Tasks

Dataset: https://www.kaggle.com/datasets/CooperUnion/cardataset

1. The dataset contains more than 10,000 rows and more than 10 columns which contains features of the car and its (MSRP) manufacturer's suggested retail price. Clean the data and analyse it making it ready for modelling.

Module 6: Supervised Machine Learning

Theoretical Assignments:

1. Explain Supervised vs. Unsupervised Learning
2. Presentation: How Regression Models are used in Business

Practical Tasks

Task 1: House Price Prediction

Dataset: https://www.kaggle.com/datasets/tusharpaul2001/house-price-prediction

The dataset contain 1460 training data points and 80 features that might help to predict the selling price of a house. Note: Save each resultant model by different name so that you can compare all of them in your final conclusion. Do as directed

1. Perform EDA
2. Build a Simple Linear Regression model to predict the Sale price of the house. Use Area as the independent variable
3. Build Multiple Linear Regression model to predict Sale price of the house.
4. Use dimensionality reduction technique PCA/LDA and build Multiple Linear Regression model to predict Sale price of the house.
5. Build a model using Lasso and Ridge regression to reduce model complexity.
6. Build an SVR model to predict Sale price of the house.
7. Build Decision Tree Regressor to predict Sale price of the house.
8. Build Random Forest Regression model to predict Sale price of the house.
9. Use GridsearchCV and RandomizedsearchCV for tuning hyperparameters and fit your model on the optimal parameters.
10. Model Selection: Evaluate and compare performance of all the models to find the best model.

Task 2: Predicting Customer Churn

Dataset: https://www.kaggle.com/datasets/thedevastator/predicting-credit-card-customer-attrition-with-m

The data is centred on customer churn, the rate at which a commercial customer will leave the commercial platform that they are currently a (paying) customer of a telecommunications company.

Do as directed

1. Perform EDA
2. Check if there is an imbalance in data. If there is an imbalance in data, resolve it.
3. Build a Logistic Regression classification model which will predict whether a customer is at risk to churn from the platform
4. Build Naive Bayes model which will predict whether a customer is at risk to churn from the platform.
5. Build a K-nearest classifier which will predict whether a customer is at risk to churn from the platform.
6. Find optimal parameters for the algorithm through GridSearchCV and build SVC model which will predict whether a customer is at risk to churn from the platform.

7. Find optimal parameters for the algorithm through GridSearchCV and build a Decision tree which will predict whether a customer is at risk to churn from the platform.
8. Find optimal parameters for the algorithm through RandomSearchCV and build a Random which will predict whether a customer is at risk to churn from the platform.
9. Model Selection: Evaluate and compare performance of all the models to find the best model.

Module 7: Unsupervised Machine Learning

Theoretical Assignments:

1. Explain Clustering algorithms and their use cases
2. Presentation: Model Evaluation techniques in Unsupervised Learning

Practical Tasks

Task 1: Customer Segmentation

Dataset: https://www.kaggle.com/datasets/vishakhdapat/customer-segmentation-clustering/data

Customer Personality Analysis is a detailed analysis of a company's ideal customers. It helps a business to better understand its customers and makes it easier for them to modify products according to the specific needs, behaviours and concerns of different types of customer segments.

Do as directed

1. Find Customer segments using K-means clustering algorithm.
2. Find Customer segments using Hierarchical clustering algorithm.
3. Compare the Clusters and conclude your analysis.

Module 8:Neural Networks & Deep Learning

Theoretical Assignments:

1. Explain the Architectures of Neural Networks
2. Presentation: Differences between CNNs, RNNs and Transformers

Practical Tasks

Task 1: Predicting Customer Churn

Dataset: https://www.kaggle.com/datasets/thedevastator/predicting-credit-card-customer-attrition-with-m

The data is centred on customer churn, the rate at which a commercial customer will leave the commercial platform that they are currently a (paying) customer of a telecommunications company. Build ANN to predict Sale price of the house.

Task 2: Digit Recognizer

Dataset: https://www.kaggle.com/datasets/hojjatk/mnist-dataset

The MNIST dataset is an inbuilt that consists of images of digits from a variety of scanned documents. Each image is a 28X28 pixel square. In this dataset 60,000 images are used to train the model and 10,000 images are used to test the model. There are 10 digits (0 to 9) or 10 classes to predict. Build CNN to predict the digit

Task 3: Sentiment Analysis

Dataset: https://www.kaggle.com/datasets/amanbarthwal/imdb-movies-data

The imdb dataset is an inbuilt that consists of reviews given by the viewer. Do as directed

a. Build RNN to Analyze the review.
b. Build LSTM to Analyze the review.
c. Build GRU to Analyze the review.
d. Model Selection: Evaluate and compare performance of all the models to find the best model suited for this application.

Task 4 : Build a text classifier model for emotion detection in text. [NLP]

Dataset: https://www.kaggle.com/datasets/kushagra3204/sentiment-and-emotion-analysis-dataset

The Dataset consists of real life emotions of the people of the US, where the data is in the csv format and you have to build a model.

Task 5: Text Classification (Average Word2vec) [NLP]

Dataset: https://www.kaggle.com/datasets/kushagra3204/sentiment-and-emotion-analysis-dataset

This is a dataset for binary sentiment classification containing substantially more data than previous benchmark datasets. We provide a set of 25,000 highly polar movie reviews for training and 25,000 for testing.

Task 5: Avengers Face Detection [CV]

Dataset: https://www.kaggle.com/datasets/rawatjitesh/avengers-face-recognition

The Dataset contains around 50 cropped face images of each avenger.

Chris Evans (Captain America)

Chris Hemsworth (Thor)

Mark Ruffalo (Hulk)

Robert DowneyJr (The Iron man)

Scarlett Johansson (Black Widow)

Module 9:Introduction to Generative AI

Theoretical Assignments:

1. Explain Generative AI vs. Traditional Machine Learning
2. Presentation: Real-World Application of Generative AI

Practical Tasks

Task 1: Generate AI-Generated Text using Gemini

- Use Google's API to generate AI-based responses.
- Deliverable: Python script generating meaningful responses.

Task 2: AI-Powered Image Generation with DALL-E

- Generate AI-created images using Gemini
- Deliverable:AI- generated images and code.

Module10-16 : Advanced LLMs, Prompt Engineering, and RAG

Theoretical Assignments:

1. Explain the key components of Transformer Model
2. Presentation: How LLMs learn from data

Practical Tasks

Task 1: Fine-Tune a Pre-Trained LLM

- Fine-tune a GPT model on a custom dataset.
- Deliverable :Model performance evaluation

Task 2:Implement a Retrieval-Augmented Generation (RAG) Pipeline

- Use Lang Chain to enhance responses with real-time data.
- Deliverable:Jupyter Notebook demonstrating the pipeline.

Module17: Capstone Project

Choose One Project:

- Sales Forecasting using Predictive Analytics
- Building a Chatbot usingGPT
- Image-based Customer Segmentation using GANs
- AutomatedDataCleaningPipelineusingPython